

Detector speed benchmark

- **Generated:** 2026-06-19 07:46
- **GPU:** NVIDIA GeForce RTX 4070
- **Method:** 1 epoch on a small symlinked subset (resolution 736); throughput extrapolated to the full 17,684-image training set.
- **MLflow:** each run is logged to the bambi - bench experiment at <http://localhost:5000> (separate from the bambi -detection training experiment).

Model	Params	Config	Subset imgs	img/s	Peak VRAM	Est. 1 epoch (17,684)	Est. 50 epochs
rf-detr-l-640	33,933,218 (33.9 M)	resolution=640, batch_size=8, grad_accum=1	500	15.02	11.29 GB	19.6 min	16.35 h
rf-detr-l	33,933,218 (33.9 M)	resolution=1024, batch_size=3, grad_accum=1	500	8.17	8.36 GB	36.1 min	30.06 h
rf-detr-s-640	32,111,170 (32.1 M)	resolution=640, batch_size=10, grad_accum=1	500	16.28	12.17 GB	18.1 min	15.09 h
rf-detr-s	32,111,170 (32.1 M)	resolution=1024, batch_size=4, grad_accum=1	500	9.07	10.94 GB	32.5 min	27.08 h
yolo26-l-640	26,299,704 (26.3 M)	imgsz=640, batch=12	500	27.97	8.87 GB	10.5 min	8.78 h
yolo26-l	26,299,704 (26.3 M)	imgsz=1024, batch=5	500	14.51	9.74 GB	20.3 min	16.93 h
yolo26-s-640	10,009,784 (10.0 M)	imgsz=640, batch=30	500	32.67	9.13 GB	9.0 min	7.52 h
yolo26-s	10,009,784 (10.0 M)	imgsz=1024, batch=12	500	26.09	9.64 GB	11.3 min	9.41 h