

Detector speed benchmark

- **Generated:** 2026-06-18 18:44
- **GPU:** NVIDIA GeForce RTX 4070
- **Method:** 1 epoch on a small symlinked subset (resolution 736); throughput extrapolated to the full 17,684-image training set.
- **MLflow:** each run is logged to the bambi - bench experiment at <http://localhost:5000> (separate from the bambi - detection training experiment).

Model	Params	Config	Subset imgs	img/s	Peak VRAM	Est. 1 epoch (17,684)	Est. 50 epochs	Status
rf-detr-l	33,933,218 (33.9 M)	resolution=736, batch_size=6, grad_accum=1	500	13.48	10.02 GB	21.9 min	18.22 h	ok
rf-detr-s	32,111,170 (32.1 M)	resolution=736, batch_size=8, grad_accum=1	500	15.35	12.08 GB	19.2 min	16.00 h	ok
yolo26-l	26,299,704 (26.3 M)	imgsz=736, batch=10	500	25.3	9.83 GB	11.6 min	9.71 h	ok
yolo26-s	10,009,784 (10.0 M)	imgsz=736, batch=24	500	44.95	9.8 GB	6.6 min	5.46 h	ok